

Why UniBrain Is FERPA Aligned by Design

FERPA is fundamentally concerned with control, disclosure, and traceability of student education records. UniBrain's ontology driven and deterministic architecture directly addresses these concerns by separating academic decision logic from language model interaction.

1. Minimization of Data Disclosure

Traditional language model centric systems transmit raw student records or transcript level data to external models as prompt context. From a FERPA perspective, this constitutes disclosure of protected education records to a third party. UniBrain avoids this entirely by computing academic outcomes internally and only passing derived, non sensitive results to the language model for explanation.

2. Deterministic and Inspectable Academic Decisions

FERPA grants students the right to inspect and challenge their academic records. UniBrain satisfies this requirement by ensuring that all academic evaluations are deterministic and reproducible. Every conclusion can be traced back to a specific requirement, slot, or policy within the ontology.

3. Clear Separation of Responsibility

UniBrain enforces a strict boundary between institutional logic and language model output. The ontology and reasoning engine are responsible for academic truth, while the language model acts solely as a narrative interface. This distinction ensures that the language model is not an academic decision maker.

4. Principle of Least Privilege

Because UniBrain's reasoning engine operates on structured data, access controls can be applied precisely. Students, advisors, and administrators can be limited to only the data necessary for their role, and all access can be logged and audited.

5. Institutional Trust and Compliance Readiness

Universities are unlikely to adopt systems that rely on opaque artificial intelligence decisions or uncontrolled data sharing. UniBrain's architecture allows institutions to retain full ownership of student records, run reasoning logic within controlled infrastructure, and replace or disable the language model if needed.

Core Principle

UniBrain does not disclose education records to language models. It discloses institutionally computed academic outcomes for explanation purposes only. This principle places UniBrain squarely within FERPA aligned system design.